

# Computer Science 360 –Operating Systems Spring 2021

## *Assignment #4*

*Due: Tuesday, April 13, 11:55 pm via push to your gitlab.csc repository*

### **Programming Platform**

For this assignment **your code must work on the Virtualbox/Vagrant configuration you provisioned for yourself in assignment #0**. You may already have access to your own Unix system (e.g., Ubuntu, Debian, macOS with MacPorts, etc.) yet I recommend you work as much as possible while with your CSC360 virtual machine. Bugs in systems programming tend to be platform-specific and something that works perfectly at home may end up crashing on a different computer-language library configuration. (We cannot give marks for submissions of which it is said “It worked on Visual Studio!”)

### **Individual work**

This assignment is to be completed by each individual student (i.e., no group work). Naturally you will want to discuss aspects of the problem with fellow students, and such discussion is encouraged. However, sharing of code is strictly forbidden. If you are still unsure about what is permitted or have other questions regarding academic integrity, please direct them as soon as possible to the instructor. (Code-similarity tools will be run on submitted work.) Any fragments of code found on the web and used in your solution **must be properly cited** where it is used (i.e., citation in the form of a comment giving the source of code).

### **Goals of this assignment**

- Complete a C program that simulates three page-replacement schemes that may be used by virtual-memory systems.

## Goal: Implement FIFO, LRU, and Second-Chance page-replacement schemes

This goal requires you to simulate the actions of a virtual-memory subsystem. Specifically you are to take a sequence of logical memory operations generated by a target process and “convert” them into physical addresses. In reality, however, you will not be running a target process but instead using a memory trace generated for you already (using an application suite from Intel called pin) and stored into a text file. You are provided with a code skeleton named `virtmem.c` which currently reads in memory operations contained in a specified trace file and calls the `resolve_address()` function for each operation. This skeleton code also finds a free frame but does not replace pages when the simulated memory is full.

Here are the first few lines from a memory trace generated by a “Hello, World!” program:

```
I: 0x7fecad0272d0
I: 0x7fecad0272d3
W: 0x7ffa4bba4d8
I: 0x7fecad02aa40
W: 0x7ffa4bba4d0
I: 0x7fecad02aa41
I: 0x7fecad02aa44
W: 0x7ffa4bba4c8
```

Each line starts with a leading “I”, “W”, or “R” (standing for “instruction read”, “memory write”, or “memory read” respectively). This single character (and a trailing colon) is followed on the line by a virtual-memory address in hexadecimal. The address must be converted into a physical address, and it is this conversion step – and much that is needed to make it happen by way of data structures and algorithms – that your code will simulate.

The skeleton code and a few trace files generated by pin are in `linux.csc.uvic.ca` within `/home/zastre/csc360/assign4`. You should copy the C source files into your Vagrantshire.

Your programming tasks within `virtmem.c` are:

- to modify `resolve_address()`;
- to modify the struct `page_table_entry` data structure (if necessary), modify `startup()` and `teardown()` functions (if necessary),
- add functions (if necessary)
- to implement a simulation the FIFO, LRU and SECONDCHANCE page-replacement algorithms.

Note that `resolve_address()` takes two parameters (the logical address and whether the operation on the address is a read or write) and returns one value (the physical address).

- The simulator is invoked on the command line when running *virtmem*.
- The algorithm to be used is indicated by a command-line argument (`--replace=fifo`, `--replace=lru` or `--replace=secondchance`).
- The size of a simulated physical frame is indicated at the command line (`--framesize=8` specifies physical frames of size 256 bytes each, i.e.,  $2^8$ ). Note that the frame size is by this forced to be a power-of-two exponent.
- The size of the simulated physical memory is indicated at the command line (`--numframes=256` specifies a simulated memory where there are 256 frames). There is no requirement that the number of frames must be a power-of-two.
- The file containing the memory trace is indicated at the command line (`--file=hello_out.txt`). These traces are provided to you, although you are free to construct your own traces. (Please see the last-page appendix to this assignment for more information on generating traces.)
- A progress-bar can be enabled via the `--progress` command-line argument.

Some functionality is already provided for you by the skeleton code. For example, it handles the processing of command line arguments. It also handles the reading of trace files, and splitting virtual memory addresses into the page-number of offset components. (Recall that these latter details depend for correct operation upon the value provided as the frame size when running the simulator.)

As an example, `ls_out.txt` contains the memory operations generated by running the Unix `ls` command on a directory in my account. Assuming you have compiled the skeleton `virtmem.c` in your account and are running it from your current directory, here is a run (without page-replacement) where the frame size is  $2^{12}$  bytes in size, where there are 256 such physical frames, and the progress bar is printed. (The scheme below is specified as FIFO, but this is ignored in the skeleton code!)

```
./virtmem --file=/home/zastre/csc360/a3/ls_out.txt --framesize=12 \  
--numframes=256 --replace=fifo --progress
```

And here is the output:

```
Progress [.....] 100%  
Memory references: 563939  
Page faults: 239  
Swap ins: 239  
Swap outs: 0
```

It so happens that the trace above was “simulated to completion”. If we change the number of frames to a smaller number:

```
./virtmem --file=/home/zastre/csc360/a3/ls_out.txt --framesize=12 \  
--numframes=100 --replace=fifo --progress
```

then here the simulation *does not* run to completion:

```
Progress [.....] 19%  
Simulator error: cannot resolve address 0x7f4038987c44 at line 108820
```

The line number message indicates that the simulator attempted to resolve the memory access indicated by line 108820 within `ls_out.txt`, but could not do so. In this case the error results from there no longer existing any more free frames (i.e., the skeleton file does not implement page replacement!).

To sum up, you are to:

- Implement a FIFO page replacement scheme and update the appropriate global variables so `output_report()` works.
- Implement an LRU page replacement scheme and update the appropriate global variable (ditto comments about `output_report()`).
- Implement an Second Change page replacement scheme and update the appropriate global variable (ditto comments about `output_report()`).
- Test your implementation with a variety of trace files, frame sizes and memory sizes.
- Make sure your operations work on 64-bit addresses (i.e., `long ints`). You’ll get weird and hard-to-debug error messages if you depend too much upon 32-bit integers (i.e., regular `ints`).

### What you must submit

- One C source-code file named `virtmem.c` containing your solution.
- One `README.md` describing (as clearly as possible) the data structures and algorithms you have used in your implementation, as well as the overall strategy of your solution.

## Evaluation

During evaluation we will use the traces provided to you, but will also use traces that will not be shared beforehand.

Our grading scheme is relatively simple.

- “A” grade: An exceptional submission demonstrating creativity and initiative. The simulator runs without any problems.
- “B” grade: A submission completing the requirements of the assignment. The simulator runs without any problems.
- “C” grade: A submission completing most of the requirements of the assignment. The simulator runs with some problems.
- “D” grade: A serious attempt at completing requirements for the assignment. The simulator runs with quite a few problems.
- “F” grade: Either no submission given, or submission represents very little work.